

structured

November 15, 2017

Contents

1	Introduction - Present legacy code	3
1.1	What is it and why do we care	3
1.2	Why is legacy code a problem	4
1.3	How have people tried to solve it	5
1.3.1	Code analysis	5
1.3.2	Manual reverse engineering	6
1.3.3	Code smells	6
1.3.4	Functional programming as a potential solution	8
1.4	Transformations to and from a legacy system	11
1.5	Biodalliance - A Legacy JavaScript Application	11
1.6	Pure web development with Purescript	12
1.7	Hypothesis and evaluation	12
2	Method	12
2.1	TODO Our code smells	12
2.1.1	TODO Long complicated functions – but only unnecessarily so.	12
2.1.2	TODO Duplicated code	12
2.1.3	TODO Primitive obsession	13
2.1.4	TODO Use of mutable state	13
2.1.5	TODO Side effects	13
2.1.6	TODO Difficult to make changes	14
2.1.7	Transformations	14
2.2	Graph Genetics Browser	14
2.2.1	TODO Specification	14
2.3	Interfacing with existing JS	14
2.3.1	FFI intro	14

2.3.2	Biodalliance	15
2.3.3	Cytoscape.js	16
2.4	Configuration	22
2.4.1	Configuring Biodalliance	22
2.4.2	BrowserConfig	26
2.4.3	Tracks	29
2.4.4	Events	29
2.5	Units	31
2.5.1	Biodalliance/The problem	31
2.5.2	Newtypes	32
2.5.3	Positions	32
2.5.4	Chromosomes	34
2.5.5	Scale	34
2.5.6	Features	34
2.6	Glyphs	35
2.6.1	Biodalliance	35
2.6.2	Glyphs in the Genetics Graph Browser	38
2.6.3	Logging glyphs	40
2.6.4	Drawing glyphs to canvas and SVG	42
2.6.5	Generating bounding boxes	43
2.6.6	Putting it all together	46
2.6.7	Limitations/Performance	48
2.7	Events	48
2.7.1	Notes	48
2.7.2	Events	48
2.7.3	Biodalliance	49
2.7.4	Cytoscape.js	49
2.7.5	Type-safe – but compile-time doesn’t make sense . . .	50
2.7.6	JSON zippers and stringy types	52
2.7.7	TrackSource and TrackSink	52
2.8	User interface	60
2.8.1	Biodalliance	61
2.8.2	Quick Halogen intro	62
2.8.3	The main application	67
3	Results	69
4	Discussion	69
5	Appendix? SVG	69

1 Introduction - Present legacy code

1.1 What is it and why do we care

Wikipedia defines "legacy system": "In computing, a legacy system is an old method, technology, computer system, or application program, 'of, relating to, or being a previous or outdated computer system.' [...] This can also imply that the system is out of date or in need of replacement."

TODO statistics on legacy systems; how pervasive are they, how old do they tend to be, how expensive is their maintenance and where do most of the costs fall

TODO sum up why it's a problem - new features, new data, security issues

TODO point out source code as one problem; not language or OS, since that's basically stable these days (i.e. they're unlikely to be replaced soon, compared to how fast things moved up until the 90s)

Following this, we define a "legacy codebase" to be the codebase of a system which is "old" and "difficult" to work with. The technology itself does not matter, only the size and complexity of the codebase itself. Likewise we do not look at dependencies, a problem solved by pure functional package managers such as Nix and Guix.

TODO What remains is "legacy code". Why is it a problem, i.e. what causes it (when the system is being designed/built) and why does it come up (when someone wants to make changes to the system)? How have people attempted to solve this problem?

"The maintenance of legacy code is a hard task since developers do not have access to details of its implementation." <http://ieeexplore.ieee.org/abstract/document/7479256/?reload=true>

1.2 Why is legacy code a problem

When someone wants to change the system. Otherwise, if it's running smoothly and does everything everyone wants of it, there is no problem.

TODO

- security fixes
- bug fixes
- changes to external dependencies
- changed features

- new features (e.g. electronic healthcare records)

WIP Why is it difficult to do these things, to change legacy code? WIP Nobody understands the code => the set of changes that are safe to make are unknown The main reason it is difficult to work with legacy code is lack of knowledge of the system and codebase. A lot of work is required to discover what changes can be done without compromising the system's current function; one must then come up with a way to fit in the desired changes (new features, bug fixes, etc.) in that possible implementation space. TODO REF

WIP There is no clear specification of the design or the system This lack of knowledge can be caused by not having a system specification, but even if there is one, it is possible (likely?) that the system has grown past it – that the specification does not actually describe the current system. A system specification does not necessarily help understand the codebase, either. TODO REF

WIP The system is only partly known; there is no way to test changes Often there are few or no tests, making it difficult, if not impossible, to be sure that a change has not caused some kind of regression. TODO REF

WIP The existing data structures etc. can be arbitrarily complex, have old/unused data/fields, be mutated at various places; thus it is difficult to know where to start when inserting new code If the system has grown organically over a longer period of time, the data structures and procedures that manipulate them have likely also grown to fit new features etc., leading to large pieces of state that are difficult to reason about. Objects may be doing something much different from their original purpose. WIP TODO REF

WIP Some new feature/solution to a new problem may involve "stuff (architecture etc.)" that the existing code is difficult to fit into, or vice versa The changes that need to be made may fundamentally be "out of phase" with the existing system. For example, it may require data that does exist in some part of the system, but the data is tangled up with other state and so on.

In short, the problem is lack of knowledge in what the system does, and how the code relates to the system and its parts. It may also be extremely difficult to know if a change made is safe.

1.3 How have people tried to solve it

Reverse engineering of legacy code exposed <https://doi.org/10.1145/225014.225045>

"Reverse engineering of large legacy software systems generally cannot meet its objectives because it cannot be cost-effective. [...] it is very costly

to “understand” legacy code sufficiently well to permit changes to be made safely, because reverse engineering of legacy code is intractable in the usual computational complexity sense."

As one of the main problems of legacy code is lack of knowledge, one of the main ways to attempt to solve it is to reverse engineer the existing system. This is usually done manually [TODO REF], but can also be done using automatic analysis of the code.

1.3.1 Code analysis

One common way is to analyze the code to find ways to modularize it, to decouple the pieces from one another. This can be done by OOP stuff

Another interesting route is finding a modularization by constructing a concept lattice based on where different global variables are used. This lattice can then be used to create descriptions on how to modularize the codebase. REF: Assessing Modular Structure of Legacy Code Based on Mathematical Concept Analysis <https://doi.org/10.1145/253228.253354>

TODO Problems with these approaches, and to automation of this in general TODO probably something about machine learning

1.3.2 Manual reverse engineering

The most common (TODO REF) solution is simply to do it by hand. This requires that programmers look at, and comprehend, the codebase, and how the codebase relates to the semantics of the system. Writing tests should also be done to give greater confidence when changes are made.

REF The most common way - Do it by hand.

It's also desirable to know what makes code more or less difficult to work with – both to give a clearer picture of what code to focus on when refactoring old systems, as well as give programmers guidelines when developing new code. In a word, we want heuristics for good vs. bad code. These are called code smells, or antipatterns.

TODO this probably needs... more First, however, we need to define what is meant by "good" or "bad" code. Simply, a good piece of code makes it clear what it does, how it relates to the system at large, and how it can be changed or reused without compromising its behavior or the behavior of the system.

TODO get Armstrong quote in here

1.3.3 Code smells

Beck & Fowler [TODO REF] provides a list of 22 code smells that have been used (extensively?) since publication. (TODO REFs <http://ieeexplore.ieee.org/abstract/document/1235447/> <https://link.springer.com/article/10.1007/s10664-006-9002-8> <https://dl.acm.org/citation.cfm?id=2629648> <http://www.sciencedirect.com/science/article/pii/S0164121215000631>)

Many of the code smells they list, as well as the solutions to them, are concerned about class-based object-oriented programming (OOP). OOP has been the primary programming paradigm for decades (TODO REF), WIP more here (but what?)

TODO is it even worth pointing out the OOP specific ones? Doesn't really apply to project

Some of the code smells:

1. General

- Duplicated code - when a piece of code appears multiple times in the codebase, it's a sign of a potential abstraction. It also makes it difficult to change things, as the change needs to be duplicated several times – which also leads to more opportunities for mistakes to be made.
- primitive obsession - Using primitive types to represent values that could be better represented by composite types or wrapped types. This can lead to values being used where they shouldn't be (e.g. providing a Number representing a pixel to a function expecting a Number actually representing the number of objects in an array). This also makes it more difficult to understand what a variable or value is to the program.
- shotgun surgery - When functionality is spread out in the codebase in such a way that making a change at one place requires making a change at many other places. Reduces code comprehension, as distant pieces of code are somehow interacting, and makes it more difficult to modify the codebase successfully.

TODO this section is probably superfluous

2. Procedural-specific

- long method - If a function or method is long, it can be difficult to understand, especially if the function interacts with implicit or global state, or performs side-effects.
- long parameter list
- switch statements
- msg chains

3. Summary What makes a code smell? Our definition of "good" from above seems to fit in well with the more general of these code smells. These code smells are generally concerned with limiting the reach of a piece of code (fewer, more organized method parameters, shorter methods, not touching other classes), as well as minimize

TODO something with SOLID maybe

TODO something like: Code whose implementation is in the neighborhood of the implementation space of semantically relevant extensions to the program.

"What we want is to increase the possible implementation space without changing the existing implementation. By transitivity, we're not doing anything in PS that cannot be done in BD. What we **are** doing, is doing this in a way such that the new implementation is closer to the intended program semantics – for some definition of "closer"."

WIP Generally, the code smells are largely concerned with procedural code. Mutability and side-effects are taken for granted; however, pure functional programming (FP) has been growing in popularity, which disallows or at least discourages functions with side-effects, and immutable data is the norm.

1.3.4 Functional programming as a potential solution

TODO As noted, OOP has been used both to fix as well as prevent legacy code issues, but has it really worked? If it has, is it the best, or even the only, way?

WIP Functional programming as (a -> b) w/o implicit state Functional programming as a paradigm focuses on functions in the mathematical sense, where functions, with some given input, always produces the same output. In purely functional programming, this concept taken to its limit, with functions not being able to perform side-effects, such as reading input from the user, or updating the user interface; more on these actions below. This is in

contrast to the imperative paradigm, which places no such limitations on functions, other than scoping. (Footnote: an imperative language could of course provide purity, but it's not exactly common, nor is it a natural part of the paradigm).

WIP Referential transparency Pure FP provides something called "referential transparency", which means that changing a piece of code to the result of running that code does not change the program(TODO REF). This makes "equational reasoning" possible(TODO REF), letting the programmer reason about parts of the program code as separate from the rest of the program. It gives the programmer confidence in what a function does.

TODO referential transparency example and counterexample

WIP Lambda calc (& purity) The purely functional paradigm can be seen as a natural extension to the lambda calculus, a model of computation invented by Alonzo Church(TODO REF), while imperative programming is closely connected to the von Neumann-style of computer on which it runs, and is similar to the idea of a Turing machine (TODO REF). Whereas a Turing machine models computation as manipulating symbols on an infinite tape of memory given a set of instructions (TODO REF), the lambda calculus models computation as function abstraction and application; the name derives from using λ to define functions. (TODO footnote about lambda/anonymous functions, maybe)

The Turing machine and lambda calculus models of computation are (as far as anyone has proven so far) equivalent, by the Church-Turing thesis. Thus any program that can be run on a theoretical Turing machine can be transformed to "run" in the lambda calculus. However, while programming languages that are built on the idea of a Turing machine are notoriously difficult to develop for (TODO REF/FOOTNOTE brainfuck, "turing-tar pits"), and are generally (TODO REF) interesting only as curiosities or for research, languages based on lambda calculus are more wide-spread. Indeed, the pure functional language Haskell is at its theoretical core a typed lambda calculus, based on System-F ω , which it compiles to as an intermediate language. (TODO REF, <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/FC>)

MAYBE TODO differentiate imperative and OOP

WIP Static types – w/ inference and powerful features Functional programming is orthogonal from type systems, but powerful type systems are closely related to pure functional languages (TODO REF). Haskell, being based on a typed lambda calculus, is a statically typed language,

WIP Pure FP as leveraging type system to ensure purity and is an example of using a powerful type system(TODO define) to capture effects that

are performed by the program – that is, letting a purely functional language express effects such as interacting with the real world (TODO like 2 refs).

Besides capturing effects, a powerful type system provides the programmer with tools to increase productivity(TODO REF), decrease bugs(TODO REF), make refactoring easier(TODO REF), and improve the programming experience in multiple ways (TODO ref & explanation, type-directed search).

NOTE this is maybe overkill, or could be moved to the PS syntax intro For example, Haskell and many(most?) other languages with similar type systems, do not have a ‘null’ value, instead encoding the possibility of lacking a value in the type system. In Haskell, the type ‘Maybe’ captures this possibility; if a function produces an ‘Int’, you can be sure that after calling the function you do indeed have an ‘Int’.

TODO maybe something about the saying "if it compiles, it works" and refactoring

WIP Category theory as 70 years of documentation in pure FP languages. (Abstractions, good ones!) As a type system gains features, the number of abstractions that can be expressed in it increases. Category Theory is a highly abstract branch of mathematics concerned with ‘categories’ consisting of ‘objects’ and ‘morphisms’ between objects. It is a rich field of research, and has over 70 years of results – and ideas and abstractions from it has been used in programming, especially pure FP. A classic example is Haskell’s use of ‘monads’, an abstraction which captures the essence of sequential computation(TODO REF). Haskell uses a monadic type for its IO system(TODO REF).

If a programmer can express their problem in the language of category theory, they gain access to 70 years of documentation concerning their problem. If the abstractions used can be expressed in the type system, the compiler can help prove that the program is correct.

(maybe footnote: for example, everything is an adjunction(TODO REF) and a monoid(TODO REF))

TODO (maybe) Partial application, currying

TODO Immutability (and how it’s getting more and more common outside FP) (follows from purity!) A lower-level part of pure FP, which has seen increased use outside of FP(TODO REF) is immutability of data. In a purely functional language, functions cannot modify data passed to them, as doing so would be to perform a side-effect – passing the same variable to two functions would not necessarily have the result expected if one function can modify the input to the other. Using data structures that are immutable by default makes reasoning about programs much easier as it removes that possible side-effect, no matter the programming paradigm.

TODO Result: referential transparency, program composition, more useful abstractions

- Lambda calculus & first-class functions
- Purity
- Types

TODO Reasons why it would work well (earlier "good code")

TODO Argue that FP is easier to comprehend, reason about

TODO Argue that it (often) decreases complexity vs. OOP code

While writing a program in a pure functional language, the programmer is encouraged by the language and environment to write code that is reusable and easy to reason about [REF Why functional programming matters]. You also get some level of program correctness, by writing code using types that correspond to the program semantics. You're able to construct transformations between data structures and compose them together – all type-checked by the compiler.

1.4 Transformations to and from a legacy system

TODO Not sure about this section. maybe rewrite to be more abstract, or cut down and use as an intro to the following two sections

WIP We extend an existing system + create a new platform blah bla To investigate using pure FP to work with, and extend, legacy systems, we will do just that. We will extend Biodalliance (BD), a JavaScript-based genome browser, with functionality to make extending it further easier, while staying backward compatible. This will be done in a minimally invasive manner, i.e. by modifying as little of BD's source code as possible.

WIP By identifying the key data structures and how sys. creates them & from what

WIP Goal: A system that feeds data to/from the legacy system, including producing modules for the legacy API, and will later subsume the legacy system Instead of modifying BD, a program that hooks into BD and communicates with it will be developed. This program will produce data compatible with BD, telling the browser how to render data in new ways, how to fetch data from new sources, etc. However, rather than be written in JavaScript, this program will be written in Purescript, a Haskell-inspired purely functional language that compiles to JS.

This program will not only be used to extend BD; instead, that will only be the first part of a new genetics browser, entirely written in Purescript.

The new browser will feature BD-compatibility, but will also embed a Cytoscape.js graph browser, as well as a genome browser written entirely in Purescript. These various components will be able to communicate with one another, and the user will be able to configure interactions between the components, to allow for new ways of exploring genetic data.

TODO Goal: A structured application that is both robust and easy to update and change

1.5 Biodalliance - A Legacy JavaScript Application

TODO describe BD

TODO describe GN2

TODO describe my earlier work on BD: adding modular renderers

TODO describe our general goal with GGB

TODO why we want BD (file format support, ease of adoption)

TODO why we don't want BD (horrible legacy code)

TODO what we do instead: TODO generating renderer modules with glyphs, TODO generating fetching modules TODO wrapping BD and controlling it from external UI

1.6 Pure web development with Purescript

WIP What Purescript is WIP Statically typed Purescript (PS) is a purely functional programming language, inspired by and in many ways similar to Haskell, sharing much of its syntax, and a powerful type system [Footnote not the **same** type system exactly] that can represent many high level abstractions. Unlike Haskell, PS compiles to JS, and can be used in the browser.

WIP Differences from Haskell WIP Good FFI - easy to wrap JS Another difference from Haskell is that PS is strictly evaluated. This lets the PS compiler output normal JS, and does not require a runtime, unlike e.g. Elm, another purely functional language that compiles to JS. PS makes use of this by providing a lightweight and powerful Foreign Function Interface (FFI), which makes it easy to interact with and wrap existing JS code.

TODO Property based testing?

1.7 Hypothesis and evaluation

TODO Hypothesis FROM Given that code smells are a cause for concern wrt. maintainability and extensibility of legacy code, find code smells – ones relevant to the Genetics Browser work we want to do – in the BD codebase,

identify the problems they imply if one were to naively try to extend the BD codebase, then identify and present a functional solution using Purescript.

2 Method

2.1 TODO Our code smells

TODO list the ones we're looking at, and why (and why not others)

2.1.1 TODO Long complicated functions – but only unnecessarily so.

maybe remove this. doesn't really apply

2.1.2 TODO Duplicated code

Duplicated code can be a sign of many potential changes and ways to refactor the code; especially an unextracted abstraction.

It's problematic because if you find you need to make a change to the "abstraction" or how it works, you need to make changes in every single piece of related dupe'd code.

It also is difficult to reason about an "abstraction" that hasn't actually been abstracted out – it is likely that each instance differs slightly, and the code provides no assistance in reasoning on a higher level; you must think the lower-level data flow, even if it's not actually semantically relevant to what the abstraction should be doing. E.g. why should a function that scrolls the view care about how the view is rendered (DOM etc.) (there are probably better examples)

FP helps deduplicate code; there are plenty of abstractions to let us compose functions and data structures to maximize reuse.

2.1.3 TODO Primitive obsession

Primitive obsession is when primitive types are used to represent parts of the system that would be better represented as types of their own. In JS, we don't really have types, so this is rampant. However, even in typed languages such as Java, it is common to e.g. represent positions as Integers or Doubles, rather than create a type that actually represents the unit corresponding to measurements of the value.

Purescript has many tools to create new types; the 'newtype' keyword is especially useful for this.

2.1.4 TODO Use of mutable state

TODO rewrite this paragraph

Mutable state is inherently difficult to reason about (citation needed). Functions and objects that refer to implicit mutable state, be it global or fields on an object, can behave differently depending on the state of the object in question; it becomes extremely difficult to reason about what a piece of code does, as it may depend, in the middle of the snippet, on some obscure field; worse, it may change some field.

2.1.5 TODO Side effects

More generally, code that performs side effects is difficult, if not impossible, to reason about. Depending on the nature and magnitude of the side effects, the effect and output of the code may change immensely, even though the code itself, and even the calling code, is the same. In short: there is no way to be certain what calling a function with side effects does – there is no way to be confident that changing it, or calling it again, is safe.

Purity solves this problem.

2.1.6 TODO Difficult to make changes

code that is tightly coupled to other parts, for no apparent reason actually this is covered by side effects, basically

2.1.7 Transformations

We want code that is free from side effects, doesn't use mutable state unless appropriate, uses types that are appropriate for the values they contain.

We also want code that is easy to reuse etc.

Transformations: From raw data to visualizations; from user input to actions; from user configuration to functions.

TODO (process metrics? if there are easy ones to get from github)

TODO Where our solutions will come from (CT etc.)

TODO How we'll go about things (piece by piece)

2.2 Graph Genetics Browser

2.2.1 TODO Specification

TODO BD

TODO Cy.js

TODO Legacy stuff
TODO New stuff

2.3 Interfacing with existing JS

The Genome Graph Browser uses BD and Cy.js, which are both written in JS. To interact with their respective APIs, we must use Purescript’s Foreign Function Interface (FFI).

2.3.1 FFI intro

Purescript’s FFI works by creating a JS source file with the same name as the PS module which is going to be interfacing with the FFI, in which we define the JS functions which we will be calling from PS. This FFI module is imported into PS using the ‘foreign import’ keywords, and providing type signatures for the values we import.

The type signatures are not validated, and there are no guarantees that the FFI functions will work – the FFI is outside the type system. Here’s an example of an FFI function which takes two values and returns their (JS-y) concatenation. In Purescript we normally have to make sure it makes sense to transform a value to a String before we print it, but Javascript has no such qualms:

```
exports.showStuff = function(a) {  
  return function(b) {  
    return function() {  
      console.log(a + b);  
    }  
  }  
}
```

Since JS doesn’t care about the types, neither do we. The type signature is polymorphic in its two arguments, and returns an effect:

```
foreign import showStuff :: forall a b. a -> b -> Eff _ Unit
```

We can also define types (and kinds, and things of other kinds) using the ‘foreign import’ syntax:

```
foreign import data JSType :: Type
```

Now, the type ‘JSType’ doesn’t have any data constructors in Purescript, so we can only create values of this type by writing an FFI function that returns it. Nor can we inspect the type without the FFI; to PS, it is entirely opaque.

2.3.2 Biodalliance

To work with Biodalliance, we define a foreign type corresponding to instances of the BD browser:

```
foreign import data Biodalliance :: Type
```

We also need an FFI function to wrap the BD browser constructor. This takes the browser constructor, another helper function, and the BD configuration as arguments:

```
foreign import initBDimpl :: forall eff.
    Fn3
    Foreign
    RenderWrapper
    BrowserConstructor
    (HTMLElement -> Eff (bd :: BD | eff) Biodalliance)
```

The output of the function is a continuation that takes an HTML element to place the BD browser in, and produces the effect to create and return the BD instance.

Biodalliance can produce events, and for GGB’s event system we need to be able to attach a handlers to parse and transmit them. We create a newtype to wrap the events from BD (to make sure we don’t use a raw event where it shouldn’t be), and an FFI function that takes a BD instance and an effectful callback, returning an effect that attaches the callback.

```
newtype BDEvent = BDEvent Json
```

```
foreign import addFeatureListenerImpl :: forall eff a.
    EffFn2 (bd :: BD | eff)
    Biodalliance
    (BDEvent -> Eff eff a)
    Unit
```

```
exports.addFeatureListenerImpl = function(bd, callback) {
    bd.addFeatureListener(function(ev, feature, hit, tier) {
```

```

        callback(feature)();
    });
};

```

2.3.3 Cytoscape.js

Like BD, we define a foreign type for the Cy.js browser instance. We also have types for the Cy.js elements, collections, and a newtype wrapper for events. Note how the CyCollection type is a type constructor:

```

foreign import data Cytoscape :: Type

-- | Cytoscape elements (Edges and Nodes)
foreign import data Element :: Type

newtype CyEvent = CyEvent Json

-- | A cytoscape collection of elements
foreign import data CyCollection :: Type -> Type

```

The Cy.js constructor is similar to BD's, except we don't need to pass any functions to it, as we have Cy.js as a dependency. We can provide a HTML element and an array of JSON objects to be used as the initial graph:

```

foreign import cytoscapeImpl :: forall eff.
    EffFn2 (cy :: CY | eff)
    (Nullable HTMLElement)
    (Nullable JArray)
    Cytoscape

```

'Nullable' is a type for dealing with 'null' in the FFI. We don't actually use 'cytoscapeImpl', instead we provide more idiomatic wrapper, so the user can use the more common and idiomatic 'Maybe':

```

cytoscape :: forall eff.
    Maybe HTMLElement
    -> Maybe JArray
    -> Eff (cy :: CY | eff) Cytoscape
cytoscape htmlEl els = runEffFn2 cytoscapeImpl (toNullable htmlEl) (toNullable els)

```

The Cytoscape.js instance can be worked with in multiple ways. Data can be added to the graph, retrieved from it, and deleted:


```

-- | Add a Collection of elements to the graph
foreign import graphAddCollectionImpl :: forall eff.
    EffFn2 (cy :: CY | eff)
    Cytoscape
    (CyCollection Element)
    Unit

graphAddCollection :: forall eff.
    Cytoscape
    -> CyCollection Element
    -> Eff (cy :: CY | eff) Unit
graphAddCollection = runEffFn2 graphAddCollectionImpl

-- | Get all elements in the graph
foreign import graphGetCollectionImpl :: forall eff.
    EffFn1 (cy :: CY | eff)
    Cytoscape
    (CyCollection Element)

graphGetCollection :: forall eff.
    Cytoscape
    -> Eff (cy :: CY | eff) (CyCollection Element)
graphGetCollection = runEffFn1 graphGetCollectionImpl

foreign import graphRemoveCollectionImpl :: forall eff.
    EffFn1 (cy :: CY | eff)
    (CyCollection Element)
    (CyCollection Element)

graphRemoveCollection :: forall eff.
    CyCollection Element
    -> Eff (cy :: CY | eff) (CyCollection Element)
graphRemoveCollection = runEffFn1 graphRemoveCollectionImpl

```

The graph layout can be controlled with the ‘runLayout’ function, which takes a ‘Layout’ value to update the Cy.js browser’s current layout:

```

-- | Apply a layout to the graph
foreign import runLayoutImpl :: forall eff.

```

```

EffFn2 (cy :: CY | eff)
Cytoscape
Layout
Unit

```

```

runLayout :: forall eff.
    Cytoscape
    -> Layout
    -> Eff (cy :: CY | eff) Unit
runLayout = runEffFn2 runLayoutImpl

```

‘Layout’ is simply a newtype wrapper over ‘String’. The native Cy.js layout function takes a ‘String’ as an argument, and with this newtype wrapper we can both easily support all the layouts supported by Cy.js – easily adding more if appropriate – while staying type-safe.

```

newtype Layout = Layout String

```

```

circle :: Layout
circle = Layout "circle"

```

1. Events

Cy.js produces events in JSON format, a newtype wrapper is used to keep things safe (and improve readability of type signatures):

```

newtype CyEvent = CyEvent Json

```

The ‘onEvent’ FFI function takes an event handler of type ‘CyEvent -> Eff a’, and a ‘String’ representing the type of event, e.g. "click" for adding a handler on click events. The function returns an effect that attaches the handler to the provided Cytoscape instance:

```

onEvent :: forall a.
    Cytoscape
    -> String
    -> (CyEvent -> Eff a)
    -> Eff Unit

```

```

exports.onEventImpl = function(cy, evs, callback) {

```

```

    cy.on(evs, function(e) {
        callback(e)();
    });
};

```

2. CyCollection

The ‘CyCollection’ type is used to work with collections of elements in the Cytoscape.js browser. As it is implemented in Purescript as a ‘foreign data import’, there is no way to create values of this type without using the FFI, e.g. with ‘graphGetCollection’. Likewise all functions that manipulate ‘CyCollection’ values must be implemented in terms of the FFI.

‘CyCollection’ is a semigroup where the binary operation is taking the union of the two ‘CyCollections’:

```

exports.union = function(a, b) {
    return a.union(b);
};

```

```

foreign import union :: forall e.
    Fn2
    (CyCollection e)
    (CyCollection e)
    (CyCollection e)

```

```

instance semigroupCyCollection :: Semigroup (CyCollection e) where
    append = runFn2 union

```

Another common interaction with a collection is extracting a subcollection. With ‘CyCollection’, we can use the ‘filter’ function for this:

```

-- | Filter a collection with a predicate
filter :: forall e.
    Predicate e
    -> CyCollection e
    -> CyCollection e

```

The FFI definition of ‘filter’ uses the Cy.js API:

```
exports.filterImpl = function(pred, coll) {
  return coll.filter(pred);
};
```

The ‘Predicate’ type is another newtype wrapper, this time of functions from the given type to Boolean. Since it’s a newtype, it can be provided to the FFI functions without unwrapping it.

```
newtype Predicate e = Predicate (e -> Boolean)
```

The Cytoscape.js API provides some basic predicates on elements, nodes, and edges. For example:

```
foreign import isNode :: Predicate Element
foreign import isEdge :: Predicate Element
```

‘Predicates’ are ‘contravariant’ in their argument, meaning they can be ‘contramapped’ over, which can be seen as the opposite of normal, ‘covariant’ functors. This is done by precomposing the ‘Predicate’ with a function ‘(a -> e)’. For example, if we have some ‘Predicate Json’, i.e. a function from JSON values to Boolean, we can contramap the ‘elementJson’ function over it, ending up with a ‘Predicate Element’. This lets us filter the Cytoscape graph with all the powerful JSON parsing tools at our disposal.

```
hasName :: Predicate Json
hasName = Predicate f
  where f json = maybe false (const true) $ json ^? _Object <<< ix "name"

elemHasName :: Predicate Element
elemHasName = elementJson >$< hasName
```

‘Predicate’ is also an instance of the ‘HeytingAlgebra’ typeclass. This lets us combine ‘Predicates’ using the normal Boolean logic combinators such as ‘&&’ and ‘||’:

```
namedNodeOrEdge :: Predicate Element
namedNodeOrEdge = (elemHasName && isNode) || isEdge
```

(a) Tests

‘CyCollection’ is unit tested to help ensure that the graph operations work as expected. For example, the edges and nodes from a graph should both be subsets of the graph:

```
let edges = filter isEdge eles
    nodes = filter isNode eles
when (not $ eles ‘contains’ edges) (fail "Graph doesn't contain its edges")
when (not $ eles ‘contains’ nodes) (fail "Graph doesn't contain its nodes")
```

Conversely, the union of the edges and nodes should be equal to the original graph, and this should be commutative:

```
(edges <> nodes) ‘shouldEqual’ eles
(nodes <> edges) ‘shouldEqual’ eles
(edges <> nodes) ‘shouldEqual’ (nodes <> edges)
```

2.4 Configuration

Software needs to be configurable. The Genetics Graph Browser (GGB) has many pieces that can and/or need to be configured by the user or system administrator. For example, what tracks are currently in the view.

There are also functions that need to be provided from the external JS, such as the Biodalliance browser constructor, and the wrapper for Purescript-defined renderers.

Configuration in standard JS solutions is not safe. A problem that can arise in JS is, if a configuration is given as a regular JS object (string, dictionary, etc.), and each configuration piece is simply assigned to the respective place in the application, there is risk of some subpiece being misconfigured, or simply missing. Worst case, the application can then crash.

2.4.1 Configuring Biodalliance

To give an idea of how configuration can take place in a legacy JS codebase, we look at BD. Many parts of BD can be configured, not just the tracks to display. All of this is provided by the user as a single JS object, and passed to the browser constructor which takes care of configuring the browser.

WIP example config A very basic browser configuration could look like this:

```
var biodalliance = new Browser({
```

```

    prefix: '../',
    fullScreen: true

    chr:      '19',
    viewStart: 30000000,
    viewEnd:  40000000,

    sources:  [{name: 'Genome',
                  twoBitURI: 'http://www.biodalliance.org/datasets/GRCm38/mm10.2b',
                  desc: 'Mouse reference genome build GRCm38',
                  tier_type: 'sequence'
                }]
  });

```

This object contains configuration of basic browser functionality (the properties ‘prefix’, which is the relative URL for icons and such data, and ‘fullScreen’ which controls how the browser itself is rendered); initial browser state (‘chr’, ‘viewStart’, ‘viewEnd’, which together define the chromosome and range of basepairs the browser displays at start); and an array of track source definitions (‘sources’), which define what data to show, and how. In this case a mouse genome sequence is the only track.

There are many more options and ways to customize the browser itself, likewise there are many different kinds of sources, and ways to configure them. All configuration is provided as JS objects, and the configuration data is used in various functions to initialize objects, from the browser itself to the tracks it displays. Since these functions are the place where the configuration options are defined, it is easy to add new configuration options; for the same reason, there is no easy way to know what a configuration option does, nor what values are legal.

We’ll have a brief look at some code smells in the configuration process, before moving on to the configuration of GGB, and how it is implemented.

1. Code Examples

- (a) Defaults and browser state BD has many features, and makes use of a lot state in its main Browser object to function. Due to this, much of the browser construction sets various pieces of state to default values, and sets others to values from the configuration, or uses some configuration data to create an initial value.

As a prime example of primitive obsession, nearly all of the fields set by the browser constructor are numbers or strings, with only a few objects. Since this is JavaScript, there is no real type checking, though some of the code makes use of basic validation. However, it is wordy, and does not provide much:

```
if (opts.viewStart !== undefined && typeof(opts.viewStart) !== 'number') {
    throw Error('viewStart must be an integer');
}

this.viewStart = opts.viewStart;
```

Verbosity alone makes it understandable that only a few of the many values set are checked like this. After various defaults are set, all other options from the configuration object are set on the browser object:

```
// 140 lines of setting default options
for (var k in opts) {
    this[k] = opts[k];
}
```

Meaning if the user has somehow set an option with the same name as one of the fields used by BD, the browser will silently use the provided value, even though it may be entirely incorrect. E.g. providing a number to a function expecting a HTML DOM element.

- (b) Entangled types (probably remove this)

Sources and styles are both tightly coupled and separate – problematic...

Neat solutions in PS include mapping ‘Source’ -> ‘Source-SansStyle’, Modelling type as (Source, Style), etc. (Isomorphic)

```
var sourcesAreEqual = sourcecompare.sourcesAreEqual;
var sourcesAreEqualModuloStyle = sourcecompare.sourcesAreEqualModuloStyle;
```

- (c) Indirection The constructor itself calls a method ‘browser.realInit()’ as it finishes. This function continues much like the constructor, preparing the browser. Finally, the method ‘browser.realInit2()’ is called.

TODO rewrite: Basically, the entire constructor, and its "sub-routines" realInit and realInit2 (those names are themselves code

smells), create ad-hoc browser elements, set a bunch of default state, some of which are just values, others are derived from other configuration or data, until the whole thing is "ready".

- (d) Event handlers WIP last sentence probably needs rewording The browser constructor also sets the various DOM event handlers used by the browser's UI. A handler can be any function that takes an event as argument, meaning it is easy to write a handler that directly takes care of e.g. updating a UI element. That is also a problem, as every handler that modifies some UI state is another possibility for interference when working with the UI.
- (e) Validation and transformations WIP This style of code is commonly seen in code throughout BD, including configuration:

```
while (sti < st.length && ry > st[sti].height && sti < (st.length - 1)) {
    ry = ry - st[sti].height - tier.padding;
    ++sti;
}
if (sti >= st.length) {
    return;
}
```

This code removes the sum of the height in pixels of the tracks from a value. It does this with external effects and state.

Stuff like this:

```
if (thisB.isDragging && rx != dragOrigin && tier.sequenceSource) {
    var a = thisB.viewStart + (rx/thisB.scale);
    var b = thisB.viewStart + (dragOrigin/thisB.scale);

    var min, max;
    if (a < b) {
        min = a|0; max = b|0;
    } else {
        min = b|0; max = a|0;
    }

    thisB.notifyRegionSelect(thisB.chr, min, max);
}

if (hit && hit.length > 0 && !thisB.isDragging) {
```



```

    if (doubleClickTimeout) {
        clearTimeout(doubleClickTimeout);
        doubleClickTimeout = null;
        thisB.featureDoubleClick(hit, rx, ry);
    } else {
        doubleClickTimeout = setTimeout(function() {
            doubleClickTimeout = null;
            thisB.notifyFeature(ev, hit[hit.length-1], hit, tier);
        }, 500);
    }
}

```

All of that to handle double clicks. Using purescript-behaviors, we could define an Event on double clicks by composition (I think). Compare to debouncing a switch with electronics vs assembly (maybe).

- (f) Code smell summary REWRITE There's no control that each part of the configuration/construction works as it should, nor is there any structure to it. These functions:
- Create and work with HTML elements
 - set default options, configuration
 - setting a whole lot of UI state, including that which is used in submenus etc.
 - Sets event handlers, which are filled with code duplication, low level handling of events, low level responses to events. Scrolling up and down with the keys is a good example: The **same** code, 80 lines long, duplicated, right after another.

2. Another approach

The solution used in GGB is to parse the configuration at the start of the program, from a raw Javascript JSON object into a Purescript type, with validation and error handling and reporting. For this I opted for purescript-foreign and purescript-argonaut, annotating all failures with error messages, which bubble up to the main configuration parser, which returns an error object or a successfully parsed configuration.

2.4.2 BrowserConfig

The type BrowserConfig represents the highest level of the GGB configuration hierarchy; it is the parsed version of the JS object provided by the user.

This is the definition:

```
newtype BrowserConfig =
  BrowserConfig { wrapRenderer :: RenderWrapper
                 , bdRenderers :: StrMap RendererInfo
                 , browser :: BrowserConstructor
                 , tracks :: TracksMap
                 , events :: Maybe
                   { bdEventSources :: Array SourceConfig
                   , cyEventSources :: Array SourceConfig
                   }
                 }
```

At this point, the specific types of the values in the record are irrelevant; the important part is that they're all Purescript types, and have been parsed and validated. The parsing is done by the `parseBrowserConfig` function, which has the following type signature:

```
parseBrowserConfig :: Foreign -> F BrowserConfig
```

NOTE: add link to source, ideally make `parseBrowserConfig` and `BrowserConfig` clickable, or add links below the script (you could generate them from Emacs tags). Also make sure this code passes the current version. Same for all others. Note that this will be your documentation too.

`parseBrowserConfig` is a function that reads a JS object containing the necessary information to start the GGB, for example which tracks are included in the view, and functions for interfacing with BD.

The pattern `'Foreign -> F a'` really says that a function named `parseBrowserConfig` is applied to `Foreign` type `F` and returns a `BrowserConfig`. This type of action is ubiquitous in the modules concerning configuration, because we use the library `'purescript-foreign'`. The type `'Foreign'` is part of Purescript and is simply anything that comes from outside Purescript, and thus must be parsed before any information can be extracted from them. `'F'` is a type synonym:

```
type F = Except (NonEmptyList ForeignError)
```

```
data ForeignError =
  JSONError String
| ErrorAtProperty String ForeignError
| ErrorAtIndex Int ForeignError
```

```
| TypeMismatch String String
| ForeignError String
```

‘Except’ is practically ‘Either’, and lets us represent and handle exceptions within the type system. In this case, the error type is a non-empty list of these possible error values. If something has gone wrong, there is at least one error message connected to it; it is simply impossible to fail a parse without providing an error message!

From the type signature, then, we see that the function name does not lie: it does attempt to parse Foreign data into BrowserConfigs, and must fail with an error otherwise. We know this, because the function does not have access to anything other than the raw configuration data, which means all the pieces of the completed BrowserConfig must be extracted from the provided configuration, or there are default values provided in the function itself.

Let’s look at one of the lines from the function definition (note: if you are new to Purescript the syntax may look strange - ignore the details, it will slowly make sense and you may appreciate the terseness in time).

```
parseBrowserConfig f = do
  browser <- f ! "browser" >>= readTaggedWithError "Function" "Error on 'browser':"

```

‘F’ is a monad, which in this case is simply an object containing state (Either a NonEmptyList or an error), so what is happening here is first an attempt to index into the "browser" property of the supplied Foreign value, followed by an attempt to read the Javascript "tag" of the value. If the tag says the value is a function, we’re happy and cast the value to the type BrowserConstructor bound to the name browser, which is later referred to when putting the eventual BrowserConfig together. If the object doesn’t have a "browser" property, or said property is not a JS function, we fail, and tell the user what went wrong.

NOTE: I would move the rest of the section to a chapter on error handling because it is actually generic:

‘readTaggedWithError’ is actually simple:

```
-- The type is:
readTaggedWithError :: forall a. String -> String -> Foreign -> F a
-- The implementation:
readTaggedWithError s e f = withExcept (append (pure $ ForeignError e)) $ unsafeReadTag

```

In words, it tries to read the tag, and if unsuccessful, appends the provided error message to the error message from `unsafeReadTagged`. Let's look at the types:

```
unsafeReadTagged :: forall a. String -> Foreign -> F a
```

```
withExcept :: forall e1 e2 a.  
              (e1 -> e2)  
            -> Except e1 a  
            -> Except e2 a
```

```
append :: forall m. Monoid m => m -> m -> m
```

In this case (of the type `F`), the use of `withExcept` would specialize to have the type:

```
withExcept :: a.  
              (a -> a)  
            -> F a -> F a
```

Another way to look at it is that `withExcept` is `map` but for the error type.

2.4.3 Tracks

Tracks configurations are different for BD tracks and Cy.js graphs, though both are provided as arrays of JSON, under different properties in the `tracks` property of the configuration object, they are treated in their respective sections.

1. Biodalliance

Tracks using BD are configured using BD source configurations; they are directly compatible with Biodalliance configurations. Because of this, there is little validation on these track configurations, as there would be no reasonable way of representing the options in Purescript, as they are spread out over the entire BD codebase. There are, for example, numerous properties which can describe from where the track will fetch data and what kind of data it is, which are logically disjoint but nevertheless technically allowed by Biodalliance (though likely with undesired results).

So, the GGB takes a hands-off approach to BD tracks, and the only validation that takes place is that a track must have a name. If it does, the JSON object is later sent, unaltered, to the Biodalliance constructor.

The Biodalliance constructor is another parameter that the configuration requires. This and the ‘wrapRenderer :: RenderWrapper’ function are required for the BD interface to function properly, and are JS functions provided by Biodalliance. (TODO note that wrapRenderer is only in a modified repo?)

2. Cytoscape.js

Cytoscape graphs are currently configured by providing a name and a URL from which to fetch the elements in JSON format.

2.4.4 Events

When a user interacts with a track, e.g. by clicking on a data point, the track can communicate the interaction to the rest of the system, including other tracks. The user can configure the structure of the events that a track produces, and what a track does when receiving an event of some specific structure, e.g. scrolling the track on receiving an event containing a position.

TODO: remove below text into the source files for documentation. You can refer to that, but I would just continue with TrackSink here.

1. Parsing the user-provided SourceConfigs

The SourceConfig and TrackSource validation is done in Either String, while the BrowserConfig parsing is done in the type Except (NonEmptyList ForeignError). To actually use these functions when parsing the user-provided configuration, we need to do a transformation like this:

```
toF :: Either String ~> Except (NonEmptyList ForeignError)
```

Fortunately, Either and Except are isomorphic - the difference between the two is only in how they handle errors, not what data they contain. There already exists a function that does part of what we need:

```
except :: forall e m a. Applicative m => Either e a -> Except e a
```

Now we need a function that brings `Either String` to `Either (NonEmptyList ForeignError)`. We can use the fact that `Either` is a bifunctor, meaning it has `lmap`:

```
lmap :: forall f a b c .
      Bifunctor f
    => (a -> b)
    -> f a c -> f b c
```

It's exactly the same as `map` on a normal functor, except it's on the left-hand type.

(TODO: idk if this is actually a good comparison) The bifunctor instance on `Either` can be seen as letting us build up a chain of actions to perform on both success and failure, a functional alternative to nested if-else statements.

The final piece we need is a way to transforming a `String` to a `(NonEmptyList ForeignError)`. Looking at the definition of the `ForeignError` type, there are several data constructors we could use. Easiest is `(ForeignError String)`, as it simply wraps a `String` and doesn't require any more information. To create the `NonEmptyList`, we exploit the fact that there is an `Applicative` instance, and use `'pure'`:

```
f :: String -> NonEmptyList ForeignError
f = pure <<< ForeignError
```

Putting it all together, we have this natural transformation:

```
eitherToF :: Either String ~> F
eitherToF = except <<< lmap (pure <<< ForeignError)
```

Now we can parse the events configuration in the `BrowserConfig` parser:

```
events <- do
  evs <- f ! "eventSources"

  bd <- evs ! "bd" >>= readArray >>= traverse parseSourceConfig
  cy <- evs ! "cy" >>= readArray >>= traverse parseSourceConfig
```

```

_ <- eitherToF $ traverse validateSourceConfig bd
_ <- eitherToF $ traverse validateSourceConfig cy

pure $ Just $ { bdEventSources: bd
               , cyEventSources: cy
               }

```

(TODO: should probably just validate in the parseSourceConfig) Note how we discard (`_ <- ...`) the results from the config validation; we only care about the validation error, since the configuration values have already been parsed.

2. Future work Typing events – types are there, just not checked (also only makes sense w/ some kinda DSL/interpreter)

2.5 Units

2.5.1 Biodalliance/The problem

It is often the case that values in programs are represented using primitive types, rather than using the fact that different units in fact can be viewed as different types. When all values are regular JS numbers, there is nothing to stop the programmer from accidentally adding a length to a weight, which is likely to lead to problems. It also becomes more difficult to comprehend what a piece of code does. TODO argue for/justify last sentence?

While they are displayed in visualizations, graphs, etc., the underlying representation is rarely anything other than a string or a number. That is, to the computer, there is no semantic difference between e.g. the position of a basepair on some chromosome, the volume of a house, or pi – all of these numbers could be used interchangeably.

WIP this is the case in Biodalliance BD uses mainly raw JS numbers and strings for representing its state and data, with a few JS objects used mainly for more complex information.

TODO examples

WIP "solutions" in JS – tagged objects One way to solve this problem in JS would be to use something like the ‘daggy’ library [TODO footnote: <https://github.com/fantasyland/daggy>], which adds tagged sum "types" to JS. The developer still needs to make sure they are used correctly, but at least the program will fail with an error if a value representing a pixel length is supplied to a function expecting a length in basepairs.

Since Purescript actually has a type checker and built-in sum types, we expect it to be easier to represent these kinds of units. In fact, Purescript lets us create new types wrapping existing types, without any runtime cost. This is done using newtypes.

2.5.2 Newtypes

Newtypes are one of the ways of creating types in Purescript. They can only have one single data constructor, and that constructor must have one single parameter, hence the intuition that they wrap an existing type. At runtime, values in a newtype are identical to values of the underlying type, which can be exploited when working with the FFI, plus there is no performance hit when using newtypes.

2.5.3 Positions

Biodalliance uses basepairs (Bp) for all position data, and stores this data as a regular Javascript Number value. It's not uncommon for data to provide its position information in megabasepairs (MBp). Obviously, treating a Bp as an MBp, or vice versa, leads to problems, but if it's just a Number being thrown around, there's no way to avoid the problem other than trusting the programmer and user to do things correctly.

As a programmer and user, I find the idea of doing so reprehensible, hence the Bp and MBp newtypes:

TODO add link to lines in Units.purs

```
newtype Bp = Bp Number
newtype MBp = MBp Number
```

To work with these, we can use pattern matching:

```
toBp :: Number -> Bp
toBp x = Bp x

fromBp :: Bp -> Number
fromBp (Bp x) = x
```

However, Purescript provides a typeclass to minimize this boilerplate, namely the 'Newtype' typeclass. The compiler derives the instance, and we can then use the generic 'wrap' and 'unwrap' functions:


```

derive instance newtypeBp :: Newtype Bp _
derive instance newtypeMBp :: Newtype MBp _

mbpToBp :: MBp -> Bp
mbpToBp x = wrap $ (unwrap x) * 1000000.0

```

Purescript also provides facilities for deriving typeclass instances for newtypes. Deriving the typeclasses used in arithmetic lets us use normal operators when working with Bp and MBp:

TODO: maybe note that most of this doesn't make very much sense, e.g. multiplying two Bp's is in fact pretty silly. Would probably be "better" to use a semigroup where $\langle \rangle$ is addition...

```

derive newtype instance eqBp :: Eq Bp
derive newtype instance ordBp :: Ord Bp
derive newtype instance fieldBp :: Field Bp
derive newtype instance euclideanRingBp :: EuclideanRing Bp
derive newtype instance commutativeRingBp :: CommutativeRing Bp
derive newtype instance semiringBp :: Semiring Bp
derive newtype instance ringBp :: Ring Bp

```

```

-- now we can do
p1 = Bp 123.0
p2 = Bp 400.0

p1 + p2 == Bp 523.0

```

TODO: needs a super basic lens primer somewhere (maybe just a footnote in the first use of it), plus readBp might not be correct The Newtype instance also gives us access to the `_Newtype` lens isomorphism:

```

_Bp :: Iso' Bp Number
_Bp = _Newtype

_MBp :: Iso' MBp Number
_MBp = _Newtype

readBp :: String -> Maybe Bp
readBp s = s ^? _Number <<< re _Bp

```

2.5.4 Chromosomes

Biodalliance represents chromosome identifiers as strings. Like with Bp, a newtype wrapper helps keep track of things:

```
newtype Chr = Chr String
derive instance newtypeChr :: Newtype Chr _
derive newtype instance eqChr :: Eq Chr
derive newtype instance ordChr :: Ord Chr
derive newtype instance showChr :: Show Chr
```

2.5.5 Scale

NOTE: This is currently only used in the Native track, however the old BD rendering stuff could/should be refactored to use the new BpPerPixel

When drawing data to the screen, we need to be able to transform between screen coordinates and the coordinates used by data. For simplicity's sake, we only care about mapping between basepairs and pixels. We represent this with another newtype wrapping Number:

```
newtype BpPerPixel = BpPerPixel Number
derive instance newtypeBpPerPixel :: Newtype BpPerPixel _

bpToPixels :: BpPerPixel -> Bp -> Number
bpToPixels (BpPerPixel s) (Bp p) = p / s

pixelsToBp :: BpPerPixel -> Number -> Bp
pixelsToBp (BpPerPixel s) p = Bp $ p * s
```

2.5.6 Features

In BD, a ‘feature’ is basically any data point. While the feature objects in BD can become arbitrarily complex as various data parsers construct them in different ways, there are only three minimal pieces of information required: what chromosome the feature is on, and what range of basepairs on the chromosome it covers.

In Purescript, we represent this type as an algebraic data type (ADT).

```
data Feature c r = Feature Chr c c r
```

For convenience, we let the compiler derive how to compare two ‘Features’ for equality and order:

```

derive instance eqFeature :: (Eq c, Eq r) => Eq (Feature c r)
derive instance ordFeature :: (Ord c, Ord r) => Ord (Feature c r)

```

There is also a smart constructor for creating ‘Features’ only with coordinates that can be transformed to basepairs.

```

feature :: forall c r. HCoordinate c => Chr -> c -> c -> r -> Feature c r
feature = Feature

```

Since ‘Feature’ has two type parameters, one for the coordinates and one for the data, and is covariant in both, we have a bifunctor instance:

```

instance bifunctorFeature :: Bifunctor Feature where
  bimap f g (Feature chr xl xr r) = Feature chr (f xl) (f xr) (g r)

```

2.6 Glyphs

A "glyph" is something that can be drawn to the browser display, as well as be exported to SVG. They are also what the user interacts with, and so have bounding boxes that are used to detect whether the user has clicked on them, to produce browser events.

2.6.1 Biodalliance

WIP Biodalliance has a bunch of classes Biodalliance has a number of Glyphs, which are classes sharing a basic interface – they have a function which draws itself to the canvas, one which produces an SVG element, as well as functions providing the bounding boxes.

WIP "higher order" glyphs The Glyphs range from basic geometric shapes such as boxes and triangles, to more complex "higher order" ones, which take other glyphs and e.g. translates them.

WIP constructor Glyphs are created using the appropriate constructor, which takes data such as position on screen, color, and so on. For example, the Glyph to create a box (rectangle) requires position, size, color, transparency, and radius (for rounding corners):

```

function BoxGlyph(x, y, width, height, fill, stroke, alpha, radius) {
  this.x = x;
  this.y = y;
  this._width = width;
  this._height = height;

```

```

    this.fill = fill;
    this.stroke = stroke;
    this._alpha = alpha;
    this._radius = radius || 0;
}

```

These fields are then used in the other methods, such as `draw()`, which draws the Glyph to a provided canvas context. All of the `draw()` methods use basic HTML5 canvas commands. A snippet of `BoxGlyph.draw()` follows; the function argument is a HTML5 canvas context to perform the drawing actions on:

```

BoxGlyph.prototype.draw = function(g) {
    var r = this._radius;
    // ...
    if (this._alpha != null) {
        g.save();
        g.globalAlpha = this._alpha;
    }

    if (this.fill) {
        g.fillStyle = this.fill;
        g.fillRect(this.x, this.y, this._width, this._height);
    }

    if (this.stroke) {
        g.strokeStyle = this.stroke;
        g.lineWidth = 0.5;
        g.strokeRect(this.x, this.y, this._width, this._height)
    }

    if (this._alpha != null) {
        g.restore();
    }
}

```

Note that the HTML5 canvas context is stateful, and commands such as `fillRect` and `strokeRect` draw shapes using the current state, which is set with e.g. the `fillStyle` and `strokeStyle` fields.

WIP `.toSVG()` using thin wrapper around DOM API BD supports exporting the browser view to SVG, which is accomplished by each Glyph

having a ‘toSVG’ method. ‘toSVG()’ returns an SVG element representing the glyph in question.

```
BoxGlyph.prototype.toSVG = function() {
    var s = makeElementNS(NS_SVG, 'rect', null,
                          {x: this.x,
                           y: this.y,
                           width: this._width,
                           height: this._height,
                           stroke: this.stroke || 'none',
                           strokeWidth: 0.5,
                           fill: this.fill || 'none'});

    if (this._alpha != null) {
        s.setAttribute('opacity', this._alpha);
    }

    return s;
}
```

WIP .min(), .max(), .height(), minY(), maxY() Constant functions ‘min()’, ‘max()’, and so on, are used to calculate the bounding boxes of glyphs, to detect whether a user has clicked on a glyph:

```
BoxGlyph.prototype.min = function() {
    return this.x;
}

BoxGlyph.prototype.max = function() {
    return this.x + this._width;
}

BoxGlyph.prototype.height = function() {
    return this.y + this._height;
}
```

WIP problems: difficult to create new glyphs, difficult to add new glyphs to rendering system The problems with this way of creating and working with glyphs largely relate to code duplication, and it being difficult to compose existing glyphs to create new ones. Having to explicitly write these various functions provide many opportunities for mistakes to sneak their way in.

WIP solution: Free monads and code generation! In GGB, we instead use a Free monad to provide a simple DSL for describing Glyphs. The DSL is then interpreted into functions for rendering it to a HTML5 canvas, as an SVG element, and bounding boxes. This vastly reduces the places where mistakes can be made, and also makes it easy to test – the canvas rendering code need only be written once and then used by all Glyphs, and it can be tested on its own.

2.6.2 Glyphs in the Genetics Graph Browser

We require some types to represent our Glyphs. First, a simple ‘Point’ type representing a point in 2D space, and the ‘GlyphF’ type which contains the commands in our Glyph DSL:

```
type Point = { x :: Number, y :: Number }
```

```
data GlyphF a =  
    Circle Point Number a  
  | Line Point Point a  
  | Rect Point Point a  
  | Stroke String a  
  | Fill String a  
  | Path (Array Point) a
```

The type parameter ‘a’ in ‘GlyphF’ is there so we can create a Functor instance. This is important, because the Free monad wraps a Functor. To reduce boilerplate, we let the compiler derive the Functor instance for GlyphF – if a type can be made into a Functor, there is only one implementation, and it is mechanical.

```
derive instance functorGlyph :: Functor GlyphF
```

The Free monad is named so because it is the

In Haskell, the definition is very simple, thanks to non-strict evaluation:

```
data Free f a = Pure a  
              | Bind f (Free f a)
```

NOTE: this is probably overkill; especially the stuff with ((,) a) Here, ‘f’ is the underlying functor, and ‘a’ is whatever value we want to return. ‘Free’ provides two value constructors; one containing only a single value

(equivalent to the ‘pure’ function in the Applicative typeclass), the other containing a value in our underlying functor, which in turn contains the next "step" in the computation in the Free monad. The Free monad can be seen as a list of commands in a DSL, where said DSL is defined entirely in the underlying functor. Another way of looking at it is as a list of functors. In fact, if the underlying functor is ‘(,) a’, that is, the type of two-element pairs where the first element is of some type ‘a’, we have a type that is isomorphic to a regular list:

```
type List a = Free ((,) a) ()

[1,2,3] ~ Bind (1,
               Bind (2,
                     Bind (3, (Pure ())))))
```

The Purescript definition of Free is more complicated, so as to be stack-safe in a strict language. However, the rest of the code is in Purescript.

The free monad constructs a list of commands, and these commands can then be interpreted into some other functor, including effectful ones. Examples will come; there is some work left before we get there. First we wrap our ‘GlyphF’ functor in ‘Free’, with a type synonym to make things cleaner:

```
type Glyph = Free GlyphF
```

Next we want to lift our ‘GlyphF’ data constructors into functions. This is done using the ‘liftF’ function, which has the following signature:

```
liftF :: forall f a. f a ~> Free f a
```

Here we use ‘liftF’ to lift two of the commands in ‘GlyphF’ to ‘Free GlyphF’, the rest are exactly analogous and elided:

```
circle :: Point -> Number -> Glyph Unit
circle p r = liftF $ Circle p r unit
```

```
stroke :: String -> Glyph Unit
stroke c = liftF $ Stroke c unit
```

```
-- and so on
```

Since it's a monad, we also can use do-notation to create glyphs, after creating some helper functions:

Now we have a number of functions which produce values in the type 'Free GlyphF'. With them, we can use Purescript's do-notation, and all the other tools that come with the Monad typeclass. Here we create a simple glyph that consists of three primitives:

```
crossedOut :: Point -> Number -> Glyph Unit
crossedOut p@{x,y} r = do
  circle p r
  line {x:x-r, y:y-r} {x:x+r, y:y+r}
  line {x:x-r, y:y+r} {x:x+r, y:y-r}
```

A Glyph, then, is simply a data structure. The interesting part lies in interpreting this data structure; or, in other words, transforming it into another data structure, especially one that performs effects. In fact, an interpreter consists of a natural transformation from the 'GlyphF' functor to some other functor.

We continue with a simple interpreter, one which transforms a 'Glyph' into a 'String', which can then be printed to console, or otherwise logged.

2.6.3 Logging glyphs

The GlyphF.Log interpreter transforms Glyphs to Strings, which we can then log to the console. To run an interpreter, we use foldFree:

```
foldFree :: forall f m. MonadRec m => (f ~> m) -> (Free f) ~> m
```

The 'MonadRec' constraint ensures that only monads supporting tail recursion can be used. Without it stack safety would be a problem. The type operator ~> denotes a natural transformation, it has the same meaning as:

```
forall a. f a -> g a
```

That is, it is parametrically polymorphic mapping between functors, and so cannot touch the contents of the functor.

For producing a String, the Writer type is a natural fit, and conveniently also has a MonadRec instance. The type of the natural transformation is then:

```
glyphLog :: GlyphF ~> Writer String
```


The definition of the function is also simple enough. For each primitive, write an appropriate string, and return the contents of the functor:

```
glyphLogN (Stroke c a) = do
  tell $ "Set stroke style to " <> c
  pure a

glyphLog (Circle p r a) = do
  tell $ "Drawing circle at (" <> show p.x <> ", " <> show p.y <>
    ") with radius " <> show r <> "."
  pure a
-- similar for the rest
```

Running the interpreter consists of applying this natural transformation to the Free GlyphF, using foldFree, and then getting the resulting String from the Writer. The function ‘showGlyph’ nearly writes itself at this point:

```
execWriter :: forall w a. Writer w a -> w

showGlyph :: forall a. Glyph a -> String
showGlyph = execWriter <<< foldFree glyphLog
```

For example, logging the process of drawing the previously defined ‘crossedOut’ glyph at the point ‘{ x: 40.0, y: 10.0 }’ with radius ‘3.0’ would produce the following output:

```
Drawing circle at (40.0, 10.0) with radius 3.0
Drawing line from (37.0, 7.0) to (43.0, 13.0)
Drawing line from (37.0, 13.0) to (43.0, 7.0)
```

2.6.4 Drawing glyphs to canvas and SVG

When drawing to canvas, we use Eff as the target for our natural transformation, and simply perform whatever canvas effects are appropriate:

```
glyphEffN :: forall eff. Context2D -> GlyphF ~> Eff (canvas :: CANVAS | eff)
glyphEffN ctx (Stroke c a) = do
  _ <- C.setStrokeStyle c ctx
  pure a
glyphEffN ctx (Circle p r a) = do
  _ <- C.beginPath ctx
```

```

_ <- C.arc ctx { x: p.x
                , y: p.y
                , r: r
                , start: 0.0
                , end: 2.0 * Math.pi
                }
_ <- C.stroke ctx
_ <- C.fill ctx
pure a
-- and so on

-- | Produce an effect to render the glyph to a canvas
renderGlyph :: forall eff. Context2D -> Glyph ~> Eff (canvas :: CANVAS | eff)
renderGlyph = foldFree <<< glyphEffN

```

SVG on the other hand uses the following type as target functor:

```
type SVG a = StateT SVGContext (Writer (Array SVGElement)) a
```

The result is a series of commands which can be used to produce the desired SVG element; this can then be rendered to the DOM:

```

interpSVGEff :: GlyphF ~> SVG
interpSVGEff (Stroke c a) = do
  SVG.setStrokeStyle c
  pure a
interpSVGEff (Circle p r a) = do
  SVG.circle p.x p.y r
  pure a
-- and so on

runSVGEff :: forall a. Glyph a -> Array SVGElement
runSVGEff = execWriter <<< (flip runStateT SVG.initialSVG) <<< foldFree interpSVGEff

-- | Render a glyph to an SVG element
renderGlyph :: forall a eff. Glyph a -> Eff ( dom :: DOM | eff ) Element
renderGlyph = SVG.renderSVG <<< runSVGEff

```

2.6.5 Generating bounding boxes

BD produces events when clicking on glyphs – which GGB make use of. To do this, BD expects four constant functions on each glyph. In Purescript,

the "bounding box" type would look like this, and could be used directly by BD:

```
type BoundingBox = { min :: Unit -> Number
                    , max :: Unit -> Number
                    , minY :: Unit -> Number
                    , maxY :: Unit -> Number }
```

When constructing glyphs in BD, each new glyph provides its own explicit bounding box. This is clearly insufficient for our purposes; instead, we make use of the fact that bounding boxes form a semigroup, and in fact also a monoid.

1. Semigroups and monoids TODO: $\langle \rangle$ can be rendered nice in latex, look that up Semigroups and monoids are concepts from abstract algebra and category theory, however they are immensely useful in pure FP, as they appear in many different areas.

A semigroup is an algebraic structure consisting of a set together with an associative binary operation. Let 'S' be the set in question and 'x', 'y', 'z' any three elements from 'S', with the binary operation ' $\langle \rangle$ '. If this following law is true, we have a semigroup:

Associativity: $(x \langle \rangle y) \langle \rangle z == x \langle \rangle (y \langle \rangle z)$

A monoid is a semigroup with one special element, an identity. The example from above is a monoid if there is an element 'e' in 'S' such that these laws apply for all elements 'x' in 'S':

Left identity: $x \langle \rangle e = x$ Right identity: $e \langle \rangle x = x$

Now we can explore how bounding boxes form a monoid.

2. Monoidal bounding boxes TODO: ref to monoids and diagrams functional pearl

The type corresponding to a glyph's position is GlyphPosition:

```
newtype GlyphPosition = GlyphPosition { min :: Number
                                       , max :: Number
                                       , minY :: Number
                                       , maxY :: Number
                                       }
```

It is a newtype wrapper over a record describing each of the four edges of the bounding box. This is a semigroup, where the binary operation produces the minimal bounding box that covers both inputs. That is, we take the minimum or maximum of the respective values, to get whichever maximizes the area covered:

```
instance semigroupGlyphPosition :: Semigroup GlyphPosition where
  append (GlyphPosition p1) (GlyphPosition p2) =
    GlyphPosition $ { min: Math.min p1.min p2.min
                      , max: Math.max p1.max p2.max
                      , minY: Math.min p1.minY p2.minY
                      , maxY: Math.max p1.maxY p2.maxY
                      }
```

Note the use of the the minimum and maximum functions from the `Math` module, and how they're really doing all the heavy lifting. For 'GlyphPosition' to be a monoid, we require an identity element. We can use the fact that the semigroup instance uses 'min' and 'max' as a hint. While there is no minimum or maximum real number, (TODO: add footnote about floating point inaccuracies... would also be better off using `Maybe`) we can cheat and use positive and negative infinity, which exist in JS. Then we have:

```
forall x. Math.min x  infinity == x
forall x. Math.max x -infinity == x
```

Now the identity 'GlyphPosition' is obvious – the minimum sides are set to positive infinity, and the maximum sides are set to negative infinity:

```
instance monoidGlyphPosition :: Monoid GlyphPosition where
  mempty = GlyphPosition { min:    infinity
                          , max:   (-infinity)
                          , minY:   infinity
                          , maxY:   (-infinity)
                          }
```

Now, with our Monoid in hand, we can write another interpreter for Glyph, using `Writer` as our monad in the natural transformation:

```

glyphPosN :: GlyphF ~> Writer GlyphPosition
glyphPosN (Stroke _ a) = pure a
glyphPosN (Circle p r a) = do
    tell $ GlyphPosition { min: p.x - (r * 1.5)
                          , max: p.x + (r * 1.5)
                          , minY: p.y - (r * 1.5)
                          , maxY: p.y + (r * 1.5)
                          }

    pure a
-- and so on

glyphToGlyphPosition :: forall a. Glyph a -> GlyphPosition
glyphToGlyphPosition = execWriter <<< foldFree glyphPosN

```

With that, we get bounding boxes for free when constructing glyphs.

3. Testing our monoid Semigroups and monoids have laws; while I'm reasonably confident in having created a Real Monoid, I prefer to have my computer make sure. To do this, I use purescript-jack, a property-based testing framework, like QuickCheck.

First, some utility functions to generate and render GlyphPositions:

TODO this is in Test.Glyph

```

type ThreeGlyphs = {l :: GlyphPosition, c :: GlyphPosition, r :: GlyphPosition}

renderGlyphs :: ThreeGlyphs -> String
renderGlyphs {l,c,r} = "{ l: " <> show l <> ", c:" <> show c <> ", r:" <> show r <

genGlyphPosition :: Gen GlyphPosition
genGlyphPosition = do
    let cf = toNumber <$> chooseInt (-10000000) (10000000)
    min <- cf
    max <- cf
    minY <- cf
    maxY <- cf
    pure $ GlyphPosition { min, max, minY, maxY }

genThreeGlyphs :: Gen ThreeGlyphs
genThreeGlyphs = do

```

```

l <- genGlyphPosition
c <- genGlyphPosition
r <- genGlyphPosition
pure $ {l, c, r}

```

The law all semigroups should abide is associativity. In Jack, we describe a Property asserting that parentheses don't matter for equality:

```

prop_semigroup :: Property
prop_semigroup =
  forAllRender renderGlyphs genThreeGlyphs \pos ->
    property $ (pos.l <> pos.c) <> pos.r == pos.l <> (pos.c <> pos.r)

```

In addition to that, monoids require that the identity element in fact be left and right identity. The Property:

```

prop_monoid :: Property
prop_monoid =
  forAll genGlyphPosition \pos ->
    property $ (pos <> mempty == pos) &&
      (mempty <> pos == pos)

```

Jack then takes care of generating GlyphPositions, ensuring that these properties hold.

TODO test output

2.6.6 Putting it all together

With these interpreters, we can create a function that produces a JS object that is compatible with BD. BD expects a glyph to have:

- a function to draw the glyph to a provided canvas
- a function to export the glyph to SVG
- functions that provide the bounding box
- optionally the relevant feature, or data point, that was used to produce the glyph

To do this, we exploit the fact that Purescript records are JS objects, by constructing a record with the appropriate properties, and transform it to a Foreign value. The main function in its entirety:

```
writeGlyph' :: forall a c r. Maybe (Feature c r) -> Glyph a -> Foreign
writeGlyph' f g = toForeign { "draw": unsafePerformEff <<< \ctx -> Canvas.renderGlyph c
                             , "min": const p.min
                             , "max": const p.max
                             , "minY": const p.minY
                             , "maxY": const p.maxY
                             , "feature": f'
                             , "toSVG": unsafePerformEff <<< \_ -> SVG.renderGlyph g
                             }
  where p = unwrap $ glyphToGlyphPosition g
        f' = toNullable $
              (\(Feature chr min max _) -> {chr, min, max}) <$> f
```

Note the use of ‘const’ to produce the constant functions that describe the bounding box, after converting the ‘Glyph’ to a ‘GlyphPosition’, and ‘unsafePerformEff’ to create functions that use the canvas and SVG interpreters to produce the output expected by BD. Since the ‘feature’ field is optional, ‘toNullable’ is used to transform an eventual ‘Nothing’ to an actual JS null, before being placed in the record.

A helper function exists for working with ‘Glyphs’ in the ‘F’ functor, which is useful when the ‘Glyphs’ were constructed in the process of parsing externally provided data. In case of failure, we produce a ‘String’ containing the errors, which is the format expected by BD:

```
writeGlyph :: forall a c r. Maybe (Feature c r) -> F (Glyph a) -> Foreign
writeGlyph f fG = case runExcept fG of
  Left errors -> toForeign $ fold $ renderForeignError <$> errors
  Right glyph -> writeGlyph' f glyph
```

In short, ‘writeGlyph’ produces data, including possible errors, in exactly the format expected by BD, while staying type safe.

NOTE these should be in discussion or something

2.6.7 Limitations/Performance

TODO inefficient – rendering tens of thousands of glyphs can be slow, each glyph setting its own stroke & fill colors, even if all glyphs look the same
NOTE: still pretty fast! 100k in 8 seconds, and (probably?) $O(n)$.

TODO cause: free monad

TODO potential solution: free applicative

2.7 Events

2.7.1 Notes

not **really** a problem in BD, however there is no checking that the features provided to listeners actually have the data expected by them, leading to a risk of runtime errors and decreased reusability

would be horrible when working with events from multiple different sources, e.g. BD and Cy.js – would end up with a bunch of nested if-else statements, searching for non-null properties. and even when you find all the properties you want, there’s no guarantee that

is the BD API also limited in what can be done? well, not really; I certainly won’t be able to do any more than featurelisteners can do (and only barely in a cleaner/more correct way)

2.7.2 Events

When working with connected data, we want to be able to interact with the data in multiple ways, to explore one data set by examining another. In the architecture of GGB, this comes down to sending events between tracks – when clicking on some data point in one track, an event containing information derived from that data point is created, and sent to other tracks that have been configured to react to those kinds of events.

In short, the system consists of four parts:

1. The browser, e.g. BD, producing raw events in response to user interaction, in whatever

format it uses

1. A track source, mapping the raw event to one used by GGB
2. A track sink, consuming GGB events into some callback that performs effects on...
3. Another browser, e.g. Cy.js.

Each part of this system should also be user-configurable, and constructed in such a way as to minimize the risk of callbacks receiving events they cannot process – we want event type safety.

We begin by looking at what events are provided by BD and Cy.js.

2.7.3 Biodalliance

BD provides several facilities for the user to add event handlers, functions that are called when the user interacts with the browser, or the browser performs some action. We are interested in only one, ‘addFeatureListener’. This function adds a handler that is called when the user clicks on a feature, i.e. on a data point in a BD track.

It receives several parameters, the DOM MouseEvent that triggered the handler, the BD feature clicked on, the track the click occurred in, and an array of other objects; for simplicity’s sake, we only look at the feature clicked on. This feature is a JS object, and can contain any information that BD parsed from the raw data, meaning two features from two different tracks can look very different, and

2.7.4 Cytoscape.js

Cy.js has a vast array of potential interactions and event handlers. We will focus on regular click-events, and thus are interested in the ‘cy.on("click")’ function, which adds on-click event handlers to the elements matching the provided selector. When no selector is provided, this matches all elements, and the handler functions similarly to the one provided to BD’s ‘addFeatureListener’.

Handlers attached with ‘cy.on()’ receive the core Cy.js graph instance, the target element (or graph) that caused the event, as well as information of what kind of event it was and when it was triggered. We’re mainly interested in the ‘target’ value, which is similar to the ‘feature’ argument in BD’s handler. Like with BD, this value contains the entire element clicked on; a big and complex JS object which can contain arbitrary data.

Both BD and Cy.js, then, produce events with unordered information of arbitrary complexity – unordered in the sense that knowledge of the data is required to extract information such as genomic position from it. Even though two pieces of data may both contain position information, there is no reason to expect the data to be found in the same place in the respective JS objects, or be of the same format. Even so, we want a

2.7.5 Type-safe – but compile-time doesn’t make sense

My first attempt, ambitious as it was, failed, and was in fact misguided from the beginning – however, it serves to illustrate the goal, and illuminate the path there. This was to represent the types of events as types in Purescript, via Purescript’s row types and polymorphic variants from `purescript-variant`.

Row types make it possible to express extensible records; they are essentially type-level maps from labels to types. For example, a record in Purescript:

```
exRec :: Record ( x :: Number, title :: String )
exRec = { x: 123.0, title: "hello" }
```

Row types can also be open, making it possible to write functions that work with any record containing at least some given fields. Here is a function that works on any record with a field named ‘title’ of type String:

```
-- { label :: Type } is sugar for Record ( label :: Type )
exRec2 :: { title :: String }
exRec2 = { title: "another record" }

titleLength :: forall r. { title :: String | r } -> Int
titleLength { title } = length title

titleLength exRec == 5
titleLength exRec2 == 14
```

variants

The use of row types is not limited to records. The package `purescript-variant` provides an implementation of polymorphic variants using row types; they are to sum types what records are to product types. For example, this function ‘`eitherOr`’ works with all possible Variants, with a default implementation for labels other than “either” and “or”. A variant with the label “either” must contain a Boolean.

```
_either = SProxy :: SProxy "either"
_or      = SProxy :: SProxy "or"
_nope    = SProxy :: SProxy "nope"

eitherOr :: forall r.
    Variant ( either :: Boolean, or :: Unit | r )
    -> String
eitherOr =
    default "neither!"
    # on _either (\b -> "either " <> show i)
    # on _or      (\_ -> "or unit")
```

```

vEither :: Variant (either :: Boolean)
vEither = inj _either true

vOr :: Variant (or :: Unit)
vOr = inj _or unit

vNope :: Variant (nope :: Maybe Int)
vNope = inj _nope (Just 543)

eitherOr vEither == "either true"
eitherOr vOr     == "or unit"
eitherOr vNope   == "neither!"

```

TODO rewrite the rest of this section The goal of using variants and rows was to provide type-safety of events. An Event would simply be a variant, and the different types of events would have different labels, and thus also different types. Producers and consumers of events would have their own rows to keep track of what they could produce and consume; as a corollary, Purescript's type checker would ensure that a consumer only receives events that it knows how to consume. In other words, a consumer could be connected to a producer if the producer's row is a subset of the consumer's row.

TODO maybe could use a somewhat more in depth description here, e.g. how events tended to be records, type Location = { chr :: Chr, pos :: Bp } etc.

This is all well and good, and my early attempts worked well. Problems arose when attempting to move from a hardcoded event flow to configuring one – this is when I realized that it doesn't make sense to have the compiler check something that needs to be configured by the user, and thus checked at runtime!

(Footnote? It may be possible using type/value-level reflection/reification, as done in Functional Pearl: implicit configurations <http://www.cs.rutgers.edu/~ccshan/prepose/prepose.pdf>)

What I actually desired was a way to express events in an easy to configure way, while also guaranteeing correctness as far as possible, with good error reporting picking up the slack where necessary.

TODO Footnote: maybe possible with reflection/reification?

2.7.6 JSON zippers and stringy types

What was needed was using a single type for all the possible events, but also providing enough data to do some kind of validation – validation on data coming from and going to the FFI, meaning it cannot be trusted whatsoever.

Since ease of configuration was another important factor, I decided to start there. JSON was the natural format to use for configuration; upon reflection, it also turned out to be a good type for events in GGB.

Having decided on JSON as the configuration format still leaves the question: what does configuring an event entail? We want the user to be able to describe what the events that come from some track look like and contain, as well as describe how the raw events are transformed into GGB events.

In most cases, this focus on the configuration format, versus the actual semantics of what the configuration data will provide, would be a sign of something being quite wrong – the format is an implementation detail.

However, in this case the format and semantics overlap. If an Event is JSON, and the configuration is given in JSON, why not use the Event as configuration? That was the inspiration that led to the current system. The user configures the event system by providing templates, or patterns, that map between raw events and the various events a track produces and consumes. It can be seen as a kind of pattern matching.

2.7.7 TrackSource and TrackSink

TODO garbage paragraph The solution consists of two types, ‘TrackSource’ and ‘TrackSink’. The former transforms events from browsers to GGB events, the latter handles received events on browser tracks.

1. TrackSource

The configuration needed for a TrackSource is a name, the JSON structure for the event to be produced, and the JSON structure of the event produced by the underlying track (e.g. Biodalliance).

For this another library will be used, instead of purescript-foreign, namely purescript-argonaut.

(a) Json decoding with Argonaut

Argonaut is a library for working with JSON in Purescript, including serializing and deserializing, as well as working with the JSON trees.

One key difference to `purescript-foreign` and its `Foreign` type, `Argonaut`'s `Json` type only corresponds to actual JSON, i.e. things that are legal in JSON formatted files. Thus, functions and other values that cannot be serialized to JSON, cannot be represented in the `Json` type.

Values of type `Json` can be decoded, or parsed, in several ways. In this case we're interested in walking arbitrary JSON trees and transforming lists of paths. Before looking at how the parsing works, here is an example of a legal `SourceConfig`:

```
{
  "eventName": "range",
  "eventTemplate": { "chr": "Chr",
                    "minPos": "Bp",
                    "maxPos": "Bp"
  },
  "rawTemplate": { "segment": "chr",
                  "min": "minPos",
                  "max": "maxPos"
  }
}
```

This defines a source that parses objects/events like this one, the JS object passed to the event handler when clicking on a feature in BD:

```
{
  // ...
  segment: "chr11",
  min: 1241230,
  max: 1270230
  // ..
}
```

Into a JS object that looks like

```
{
  chr: "chr11",
  minPos: 1241230,
  maxPos: 1270230
}
```

This is useful if several tracks produce events with the same data but in objects that look different; the consumer of the event will only see events of this last format. The templates provided can be of arbitrary depth and complexity; the only rule is that each leaf is a key, and all properties be strings (i.e. no arrays). There is some validation too, detailed later.

‘eventTemplate’ and ‘rawTemplate’ are both whole structures which we’re interested in. For each leaf in the eventTemplate (including its property name), we create a path to where the corresponding value will be placed in the finished event. Similarly, we need to grab the path to each leaf in the rawTemplate, so we know how to grab the value we need in the finished event, from the provided raw event.

Fortunately, Argonaut provides functions for dealing with exactly this. First, the JCursor type describes a path to a point in a JSON tree:

```
data JCursor =
  JIndex Int JCursor
  JField String JCursor
  JCursorTop
```

It can be seen as a list of accessors. If we have an object in JS:

```
let thing = { x: [{a: 0},
                  {b: {c: true}}
                ]};
```

We can grab the value at ‘c’ with

```
let cIs = thing.x[1].b.c;
```

With JCursor, this accessor chain ‘x¹.b.c’ would look like:

```
(JField "x"
 (JIndex 1
  (JField "b"
   (JField "c" JCursorTop))))
```

It’s not pretty when printed like this, but fortunately not much direct manipulation will be needed. We create these JCursors

¹DEFINITION NOT FOUND.

from a JSON structure like the templates above with the function `toPrims`:

```
toPrims :: Json -> List (Tuple JCursor JsonPrim)
```

The type `JsonPrim` can be viewed as exactly what it sounds like – it represents the legal JSON primitives: null, booleans, numbers, strings. In this case we only care that they are strings.

This function walks through a given JSON object, and produces a list of each leaf paired to the `JCursor` describing how to get to it. That is, it does exactly what we want to do with the `rawTemplate` from earlier.

With the `eventTemplate` we don't want to pick out the leaf, but the label of the leaf. In this case we do need to step into the `JCursor` structure, but only a single step, after reversing it:

```
insideOut :: JCursor -> JCursor
```

```
eventName <- case insideOut cursor of
    JField s _ -> Just s
    _          -> Nothing
```

The function 'insideOut' does what expected and reverses the path through the tree. We then match on the now first label, and save it as the name. If it was an array, we fail with a `Nothing`.

Argonaut, especially the functions concerning `JCursor`, largely uses the `Maybe` type. This is fine for the most part, but as this will be used in configuration, and thus needs to tell the user what has gone wrong if the provided configuration is faulty, it's not enough.

A more appropriate type would be `Either String`, which allows for failure to come with an error message. To "lift" the functions using `Maybe` into `Either String`. See `source code` for an example. To provide the user with additional help when configuring, the source configurations are validated to make sure the given JSON structures are legal, or "match". Given some value that we want to have in the finished event, and all of the values we know we can get from the raw event, if we can't find the first value among the latter, something's wrong.

The implementation is simple. The Cursors here are grabbed from the result of `toPrims` above; the `JCursors` themselves are unaltered.

```
-- This is just a nicer version of Tuple JCursor String
type Cursor = { cursor :: JCursor
                 , name  :: String
                 }
```

```
type RawCursor = Cursor
type ValueCursor = Cursor
```

```
validateTemplate :: Array RawCursor
                 -> Array ValueCursor
                 -> Either String ValueCursor
validateTemplate rcs vc =
  if any (\rc -> vc.name == rc.name) rcs
  then pure vc
  else throwError $ "Event property " <> vc.name <> " is not in raw template"
```

In words, if one of the many raw event cursors has the same name as the given value cursor, it's good, otherwise throw an error. To increase this to validate the array of cursors defining a finished event, we can make use of `Either's` `Applicative` instance, and `traverse`:

```
-- specialized to Either String and Array
traverse :: forall a b.
           (a -> Either String b)
         -> Array a
         -> Either String (Array b)

validateTemplates :: Array RawCursor
                 -> Array ValueCursor
                 -> Either String (Array ValueCursor)
validateTemplates rcs = traverse (validateTemplate rcs)
```

The function tries to validate all given templates, and returns the first failure if there are any. Validation of a collection of things for free!

2. TrackSink

TrackSinks are configured by providing an event name and a callback. On the PS side, these are type-safe, but there is no way to ensure that functions passed from Javascript to Purescript are type-safe. BD and Cy.js TrackSinks, respectively, should have the following types:

```
newtype TrackSink a = TrackSink (StrMap (Json -> a))

type BDTrackSink = TrackSink (Biodalliance -> Eff Unit)
type CyTrackSink = TrackSink (Cytoscape -> Eff Unit)
```

These are the "expanded" types, for clarity. Note that they are extremely similar; the only difference is what type of browser they work on:

```
BDTrackSink = TrackSink (StrMap (Json -> Biodalliance -> Eff Unit)
CyTrackSink = TrackSink (StrMap (Json -> Cytoscape -> Eff Unit)
```

The event name is used to place the function in the correct index of the StrMap. The callback uses currying to take both the event (as JSON) and the respective browser instance, to be used e.g. when scrolling the Biodalliance view to an event.

The following JS code defines a Biodalliance TrackSink (and is correctly typed):

```
var bdConsumeLoc = function(json) {
  return function(bd) {
    return function() {
      bd.setLocation(json.chr,
                     json.pos - 1000000.0,
                     json.pos + 1000000.0);
    };
  };
};

var bdTrackSinkConfig = [ { eventName: "location",
                           eventFun: bdConsumeLoc
                         }
];
```

3. Running TrackSources and TrackSinks

For TrackSource and TrackSink to be usable we need to be able to create them from the provided configurations, and provide functions for applying them to events as appropriate.

(a) TrackSource

To create a TrackSource, the provided templates are parsed and validated. Since a TrackSource is a list of parsers, if the SourceConfig is correct, a function from raw events to parsed events is returned, wrapped in a singleton list and the TrackSource type:

```
makeTrackSource :: SourceConfig -> Either String (TrackSource Event)
makeTrackSource sc = do
  rawTemplates <- parseRawTemplateConfig sc.rawTemplate
  eventTemplates <- validateTemplates rawTemplates
  =<< parseTemplateConfig sc.eventTemplate

  pure $ TrackSource $ singleton $ \rawEvent -> do
    vals <- parseRawEvent rawTemplates rawEvent
    evData <- fillTemplate eventTemplates vals
    pure $ { name: sc.eventName, evData }
```

To extend the above function to work on a collection of configuration objects, we use function composition to first attempt to use each provided configuration to create a TrackSource, followed by combining the list of parsers into a single one:

```
makeTrackSources :: Array SourceConfig -> Either String (TrackSource Event)
makeTrackSources = map fold <<< traverse makeTrackSource
```

First ‘traverse’ is used to try to create the TrackSources, which returns an array of ‘TrackSource Event’ if all were legal, or an error if something went wrong. Next, ‘map’ is used to apply a function to the ‘Right’ side of the ‘Either’ from the use of ‘traverse’, and the applied function is ‘fold’, which concatenates a collection of values of some monoid into a single value – the monoid in question is TrackSource.

This is not the only reasonable way of defining this function – one may very well want to collect the error messages while returning the successes. As ‘makeTrackSources’ demonstrates, not much code is needed to compose functions to provide the validation logic

that is desired, and there is nothing unique about this function; all that is required is swapping out some of the functions.

Finally, a way to use a `TrackSource` is required. The following function uses a `TrackSource` to parse a provided JSON value:

```
runTrackSource :: TrackSource Event -> Json -> Array Event
runTrackSource (TrackSource ts) raw = filterMap (_ $ raw) ts
```

It works by applying each function in the array wrapped by `TrackSource` to the provided value, filtering out the ‘Nothing’s and returning an array of successfully parsed ‘Events’.

(b) `TrackSink`

A `TrackSink` is a map from event names to a function that handles the event, so to make one we create a singleton map from the provided event name to the provided function, and wrap it in the `TrackSink` type:

```
makeTrackSink :: SinkConfig ~> TrackSink
makeTrackSink sc = TrackSink $ StrMap.singleton sc.eventName sc.eventFun
```

Using a collection of ‘SinkConfigs’ to produce a single ‘TrackSink’ is not in itself complicated, but we do some validation, namely ensuring that there are not multiple handlers for a given event:

```
makeTrackSinks :: forall a.
    Array (SinkConfig a)
    -> Either String (TrackSink a)
makeTrackSinks scs = do
    let count = StrMap.fromFoldableWith (+) $ map (\c -> Tuple c.eventName 1) scs
        overlapping = StrMap.filter (_ > 1) count

    when (not StrMap.isEmpty overlapping)
        let error = foldMap (append "\n" <<< show) $ StrMap.keys overlapping
        in throwError $ "Overlapping tracksinks!\n" <> error

    pure $ foldMap makeTrackSink scs
```

In this case, we use ‘foldMap’ to map the ‘makeTrackSink’ function over the provided configurations, and then use the ‘TrackSink’ monoid instance to combine them – similar to ‘fold << traverse’ in the case of ‘TrackSource’.

To use a ‘TrackSink’, we see if a handler for the provided event exists. If it does, we apply it to the contents of the event:

```
runTrackSink :: forall a. TrackSink a -> Event -> Maybe a
runTrackSink (TrackSink sink) event = do
  f <- StrMap.lookup event.name sink
  pure $ f event.evData
```

However, since ‘TrackSinks’ are intended to perform effects, a helper function for that is useful. In particular, the following function creates a "thread" (TODO footnote on JS singlethreaded) that reads events from a provided ‘BusRW’ (TODO define/refer to ‘BusRW’ intro), running effectful functions from the provided ‘TrackSink’ if the received event has a handler:

```
forkTrackSink :: forall env.
  TrackSink (env -> Eff Unit)
-> env
-> BusRW Event
-> Aff Canceler
forkTrackSink sink env bus = forkAff $ forever do
  event <- Bus.read bus

  case runTrackSink sink event of
    Nothing -> pure unit
    Just f   -> liftEff $ f env
```

2.8 User interface

The main function of GGB’s user interface is to tie the browser tracks – BD and Cy.js – together. It also creates and to some extent manages the JS browser track instances, and renders the HTML for the entire UI.

2.8.1 Biodalliance

WIP BD intro Biodalliance has a full-featured UI for exploring genomic data chromosome-wise, adding and removing currently displayed tracks, configuring browser options, and exporting the current browser view to various formats. BD accomplishes this by creating and working with DOM elements and the HTML5 canvas API, and setting handlers on DOM events such as clicking and dragging the browser view, or pressing the arrow keys to scroll.

WIP DOM actions Because BD does not use any abstracting library for dealing with the DOM, and likely because BD has grown features organically over time, the code for updating the UI is interleaved with other code, including event handlers, fetching data for a track, and more. BD also programmatically sets various CSS properties on UI elements, and uses the web browser's computed stylesheet to figure what manipulations are necessary.

TODO example

TODO Events NOTE: mainly covered in events.org

TODO example

WIP Problems In short, BD's UI uses plenty of global state, and is highly complex and spread out over the codebase. Adding a UI element would require finding a place in the DOM where it would fit – both in screen estate as well as in styling – and somehow suture it into the code while making sure that the existing UI elements are not negatively affected by this sudden new element, plus that the other UI elements and functionality do not interact with the element in some undesired manner.

Another problem, that could arise when adding some feature, not necessarily modifying the UI itself, is the risk of the interface ending up in an inconsistent state. With all the global state that is used, both in the DOM and in the BD browser itself, it is difficult to know what changes can be made. One cannot even call a function which performs some action when a button is clicked, without risking that the function itself toggles some state.

WIP How we do it in PS In Purescript, we do not juggle DOM elements and events. Instead, we use Halogen, from `purescript-halogen`, a type-safe UI library, similar in purpose to React. Event passing between tracks is taken care of by `purescript-aff-bus` and threads from `purescript-aff`, while DOM events are handled by Halogen.

Using these tools, we can construct a potentially complex UI, with some, albeit not absolute, confidence that the UI will not move to an inconsistent state. Halogen also provides a DSL for declaratively constructing the DOM of our application. Naturally, there is no implicit global state to be concerned about.

2.8.2 Quick Halogen intro

Halogen is a component-based UI library, using a virtual DOM implementation to render HTML to the browser. A component is a value of the (fairly complicated) type `Component` (removed constraints etc. for clarity):

```
component :: Component renderer query state message monad
```

(1) (2) (3) (4) (5) (6)

The type ‘Component (1)’ takes five type parameters. The first, ‘renderer (2)’ is the type used to render the page, we use a HTML renderer. Next is ‘query (3)’, which is filled with our query algebra, to be explained later; in short it is the set of commands the component can respond to. ‘state (4)’ is the type of the state kept by the component. We don’t have any, so we set it to ‘Unit’. ‘message (5)’ is the type of messages we can produce, which we can send to other parts of the program. Finally, ‘monad (6)’ is the type in which all effects produced by the component will happen. In our case, it’s the ‘Aff’ monad for asynchronous effects – it could also be a monad transformer stack, or some free monad.

1. Query algebras

The "Query algebra" is the type describing the possible actions we can query the component to perform. The type is not complicated; in GGB we have:

```
data Query a
  = CreateBD (forall eff. HTMLElement -> Eff (bd :: BD | eff) Biodalliance) a
  | PropagateMessage Message a
  | BDScroll Bp a
  | BDJump Chr Bp Bp a
  | CreateCy String a
  | ResetCy a
```

From top to bottom, we can ask it to ‘CreateBD’, providing a function that creates a Biodalliance instance given a HTML element to place it in; we can propagate messages from the child components; we can scroll and jump the BD instance; and we can create and reset the Cy.js instance. That’s what the queries look like, but we also need to define an ‘eval’ function. This maps Query to Halogen commands, which are also defined by a functor type – the function is a natural transformation from our Query DSL to the Halogen DSL (a free monad).

```
eval :: Query ~> HalogenM state query childQuery childSlot message monad
```

The type parameters of ‘HalogenM’ are the same as those of ‘Component’, adding a ‘childQuery’ type, the Query type of values which this

component can use to communicate with its children, and ‘childSlot’, the type which is used to index into the child components. For the main GGB component they are:

```
type ChildSlot = Either2 UIBD.Slot UICy.Slot
```

```
type ChildQuery = Coproduct2 UIBD.Query UICy.Query
```

‘ChildSlot’ is a coproduct of the two child Slot *types* (Either2) of the child components; we can query the BD slot or the Cy.js slot at once, not both. ‘Either2’ is a generalization of ‘Either’ to a variable number of types, a convenience that makes it easy to change the number of slots, without more work than a type synonym. ‘ChildQuery’ is a coproduct of the two child Query *functors* (Coproduct2).

```
data Either a b = Left a | Right b
```

```
data Coproduct f g a = Coproduct (Either (f a) (g a))
```

```
-- can be viewed as (pseudocode):
```

```
data Coproduct f g a = Coproduct (Left (f a)) | (Right (g a))
```

```
type ChildQuery a = Either (UIBD.Query a) (UICy.Query a)
```

– TODO: not sure, but it may even be impossible to do this; may not compile – (certainly doesn’t compile when applied to Halogen)

We can’t use normal ‘Either’ for ChildQuery, as we wouldn’t be able to be parametric over the ‘a’ type in both child queries. If we were to map a function ‘UICy.Query (a -> b)’ on the Right component of the Either ChildQuery, we’d end up with the type ‘Either (UIBD.Query a) (UICy.Query b)’, which obviously is not congruent to ‘ChildQuery a’.

Writing the function is simple enough. We pattern match on the input Query, and produce effects in the HalogenM type. Creating BD is done by querying the BD child using its respective slot and a ChildPath – a type describing a path to the child component, and providing an action to tell the child component to perform.

```
eval = case _ of
```

```
CreateBD bd next -> do
  _ <- H.query' CP.cp1 UIBD.Slot $ H.action (UIBD.Initialize bd)
  pure next
```

‘H.action’ is a Halogen function mapping ChildQuery constructors to concrete actions, by simply applying the ‘Unit’ type to it.

```
type Action f = Unit -> f Unit
action :: forall f. Action f -> f Unit
action f = f unit
```

Finally, we return the next command. Next is ‘PropagateMessage’, which receives a Message (sent from the function handling messages from the children):

```
data Message
  = BDInstance Biodalliance
  | CyInstance Cytoscape
```

Depending on which message it is, we print a log message, and then use ‘H.raise’ to send the message out from Halogen to subscribers elsewhere in the app (more on that later).

```
PropagateMessage msg next -> do
  case msg of
    BDInstance _ -> liftEff $ log "propagating BD"
    CyInstance _ -> liftEff $ log "propagating Cy"
  H.raise msg
  pure next
```

The rest are simple queries to the respective child component, practically the same as ‘CreateBD’:

```
BDSroll dist next -> do
  _ <- H.query' CP.cp1 UIBD.Slot $ H.action (UIBD.Scroll dist)
  pure next
BDJump chr xl xr next -> do
  _ <- H.query' CP.cp1 UIBD.Slot $ H.action (UIBD.Jump chr xl xr)
  pure next
```



```

CreateCy div next -> do
  _ <- H.query' CP.cp2 UICy.Slot $ H.action (UICy.Initialize div)
  pure next
ResetCy next -> do
  _ <- H.query' CP.cp2 UICy.Slot $ H.action UICy.Reset
  pure next

```

2. Rendering Next is rendering the component. This is done by providing a function from the component 'state' to a description of the DSL used by the 'renderer' type. In our case, we render to 'HTML', and so use the type 'ParentHTML', which contains all the types required to interact with the children.

```
render :: State -> ParentHTML query childQuery childSlot m
```

The function itself is simple, we use Arrays and some functions to describe the HTML tree, a simplified version follows:

```

render _ =
  HH.div_
    [ HH.button
      [ HE.onClick (HE.input_ (BDScroll (Bp (-1000000.0)))) ]
      [ HH.text "Scroll left 1MBp" ]

    , HH.div
      [] [HH.slot' CP.cp1 UIBD.Slot UIBD.component unit handleBDMMessage]
    ]

```

This produces a button with the text "Scroll left 1MBp", and clicking on it sends a query to 'eval' to scroll the BD view 1 MBp to the left; as well as a div with the BD child component. Adding the child component here is how we create the component, so we must also provide a handler in the parent for messages from the child, namely 'handleBDMMessage'.

3. Messages A component can send messages to its parent, or the rest of the application in the case of the top-level component. These are the messages the BD and Cy.js components can produce, respectively:

```
data UIBD.Message
  = SendBD Biodalliance
```

```
data UICy.Output
  = SendCy Cytoscape
```

The main component can produce these:

```
data Message
  = BDInstance Biodalliance
  | CyInstance Cytoscape
```

Note that the main container uses its own messages to propagate the children components; message passing is limited by Halogen, and anything more complex than this should be done on another channel (which is what GGB does with events).

The messages from the BD and Cy.js components are handled by the functions ‘handleBDMessage’ and ‘handleCyMessage’:

```
handleBDMessage :: UIBD.Message -> Maybe (Query Unit)
handleBDMessage (UIBD.SendBD bd) = Just $ H.action $ PropagateMessage (BDInstance bd)

handleCyMessage :: UICy.Output -> Maybe (Query Unit)
handleCyMessage (UICy.SendCy cy) = Just $ H.action $ PropagateMessage (CyInstance cy)
```

Note that these produce Queries on the main component. We want to send the messages containing the references to the instances out from the component to the outside application, hence creating a PropagateMessage query wrapping the reference. As seen in ‘eval’ above, this in turn calls ‘H.raise’ on the message, sending it to the outside world.

4. Creating the component These functions, including one to produce the initial state (simply ‘const unit’) are all put together and provided to the ‘parentComponent’ function, producing the Component itself. This can then be provided to Halogen’s ‘runUI’ function, along with the initial state and an HTML element to be placed in, to create and run the Halogen component.

First, however, we need a ‘main’ function application to run.

2.8.3 The main application

‘main’ is the function which will be called by the user to run the browser. It takes a ‘Foreign’ object – the one to parse into a browser configuration – and then does some stuff with Eff (e.g. be a genetics browser):

- TODO: remove row blank when compiling with 0.12 – TODO: explain runHalogenAff

```
main :: Foreign -> Eff _ Unit
main fConfig = HA.runHalogenAff do
```

First we attempt to parse the provided configuration, logging all errors to config on failure, otherwise continuing:

```
case runExcept $ parseBrowserConfig fConfig of
  Left e -> liftEff $ do
    log "Invalid browser configuration:"
    sequence_ $ log <<< renderForeignError <$> e

  Right (BrowserConfig config) -> do
```

With a validated config, we can create the track/graph configs, and create the function which will later be used to create Biodalliance:

```
let {bdTracks, cyGraphs} = validateConfigs config.tracks

    opts' = sources := bdTracks.results <>
            renderers := config.bdRenderers

liftEff $ log $ "BDTrack errors: " <> foldMap ((<>) ", ") bdTracks.errors
liftEff $ log $ "CyGraph errors: " <> foldMap ((<>) ", ") cyGraphs.errors

let mkBd :: (forall eff. HTMLElement -> Eff (bd :: BD | eff) Biodalliance)
    mkBd = initBD opts' config.wrapRenderer config.browser
```

After picking the element to run in, we create the Halogen component, and create the Buses to be used by the events system. Note that we bind the value of ‘runUI’ to ‘io’:

```
io <- runUI component unit el'

busFromBD <- Bus.make
busFromCy <- Bus.make
```

‘io’ can be used to subscribe to messages sent from the main component, as well as send queries to it, which we do momentarily. First, we use the provided TrackSink and TrackSource configurations to create the BD TrackSink and TrackSource:

```
let bdTrackSink = makeTrackSinks <<< _.bdEventSinks =<<
    note "No BD event sinks configured" (config.events)
    bdTrackSource = makeTrackSources <<< _.bdEventSources =<<
        note "No BD event sources configured" (config.events)
```

We create the respective values, adding an error message if something went wrong.

Finally, we attach a callback to the Halogen component to listen for the reference to the BD instance, sent by the BD component upon creation. We then use the TrackSink and TrackSource configurations to hook BD up to the event system. Finally, we ask the main component to create the BD instance:

```
io.subscribe $ CR.consumer $ case _ of
    BDInstance bd -> do

        case bdTrackSink of
            Left err -> liftEff $ log "No BD TrackSink!"
            Right ts -> forkTrackSink ts bd busFromCy *> pure unit

        liftEff $ case bdTrackSource of
            Left err -> log err
            Right ts -> subscribeBDEvents ts bd busFromBD

        --TODO remove BDBRef? debug stuff...
        liftEff $ setBDBRef bd
        pure Nothing

    _ -> pure $ Just unit

io.query $ H.action (CreateBD mkBd)
```

If the ‘TrackSink’ was correctly configured, ‘forkTrackSink’ is used to pipe events from the Cytoscape.js instance to the handler defined by said ‘TrackSink’. We don’t care about being able to kill the "thread" using the

‘Canceler’, so we throw away the result with ‘*> pure unit’. Similarly, the ‘TrackSource’ is used with the helper function ‘subscribeBDEvents’, defined thusly:

```
subscribeBDEvents :: forall r.
    (TrackSource Event)
    -> Biodalliance
    -> BusRW Event
    -> Eff _ Unit
subscribeBDEvents h bd bus =
    Biodalliance.addFeatureListener bd $ \obj -> do
        let evs = runTrackSource h (unwrap obj)
        traverse_ (\x -> Aff.launchAff $ Bus.write x bus) evs
```

It adds an event listener to the provided BD browser instance and writes the successful parses to the provided Bus.

The Cytoscape.js code is analogous.

3 Results

4 Discussion

5 Appendix? SVG